

Name: Aditya Tanaji Gawade Roll: 75 Div: TA3

```
import pandas as pd

# A tiny hand-made dataset to illustrate the concepts above
student_data = {
    "RollNo": [1, 2, 3, 4, 5],
    "Name": ["Aarav", "Diya", "Kabir", "Meera", "Ishan"],
    "Age": [19, 20, 19, 21, 20],
    "Gender": ["M", "F", "M", "F", "M"],
    "CGPA": [8.4, 9.1, 7.8, 8.9, 7.2],
    "Placed": ["Yes", "Yes", "No", "Yes", "No"],
}

student_df = pd.DataFrame(student_data)
student_df  # displaying the DataFrame
```

	RollNo	Name	Age	Gender	CGPA	Placed
0	1	Aarav	19	M	8.4	Yes
1	2	Diya	20	F	9.1	Yes
2	3	Kabir	19	M	7.8	No
3	4	Meera	21	F	8.9	Yes
4	5	Ishan	20	M	7.2	No

```
# Core libraries
import numpy as np
import pandas as pd

# Visualization
import matplotlib.pyplot as plt
import seaborn as sns

# Preprocessing (used later for scaling & encoding)
from sklearn.preprocessing import (
    StandardScaler, MinMaxScaler, RobustScaler, Normalizer,
    LabelEncoder, OneHotEncoder, OrdinalEncoder
)

# Display / plotting settings
%matplotlib inline
sns.set_theme(style="whitegrid")
plt.rcParams["figure.figsize"] = (7, 4)
pd.set_option("display.max_columns", None)

print("All libraries imported successfully!")
```

All libraries imported successfully!

```
# Load all datasets we will use in this notebook
titanic = sns.load_dataset("titanic")
iris = sns.load_dataset("iris")
tips = sns.load_dataset("tips")
penguins = sns.load_dataset("penguins")

print("Titanic :", titanic.shape)
print("Iris :", iris.shape)
print("Tips :", tips.shape)
print("Penguins:", penguins.shape)
```

```
Titanic : (891, 15)
Iris : (150, 5)
Tips : (244, 7)
Penguins: (344, 7)
```

```
# California Housing – a larger, purely numerical dataset (fetched from scikit-learn)
from sklearn.datasets import fetch_california_housing

try:
    housing_bunch = fetch_california_housing(as_frame=True)
    housing = housing_bunch.frame
except Exception as e:
    # Fallback: if there is no internet access, build a similar synthetic dataset
    print("Could not download California Housing, generating a similar synthetic dataset instead.")
    from sklearn.datasets import make_regression
    X, y = make_regression(n_samples=2000, n_features=8, noise=15, random_state=42)
```

```

cols = ["MedInc", "HouseAge", "AveRooms", "AveBedrms",
        "Population", "AveOccup", "Latitude", "Longitude"]
housing = pd.DataFrame(X, columns=cols)
housing["MedHouseVal"] = y

print("California Housing:", housing.shape)
housing.head()

```

California Housing: (20640, 9)

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude	Longitude	MedHouseVal
0	8.3252	41.0	6.984127	1.023810	322.0	2.555556	37.88	-122.23	4.526
1	8.3014	21.0	6.238137	0.971880	2401.0	2.109842	37.86	-122.22	3.585
2	7.2574	52.0	8.288136	1.073446	496.0	2.802260	37.85	-122.24	3.521
3	5.6431	52.0	5.817352	1.073059	558.0	2.547945	37.85	-122.25	3.413
4	3.8462	52.0	6.281853	1.081081	565.0	2.181467	37.85	-122.25	3.422

```

import seaborn as sns
sns.get_dataset_names()

```

```

['anagrams',
 'anscombe',
 'attention',
 'brain_networks',
 'car_crashes',
 'diamonds',
 'dots',
 'dowjones',
 'exercise',
 'flights',
 'fmri',
 'geyser',
 'glue',
 'healthexp',
 'iris',
 'mpg',
 'penguins',
 'planets',
 'seaice',
 'taxis',
 'tips',
 'titanic']

```

```

import plotly.express as px
dir(px.data)

```

```

['_all_',
 '__builtins__',
 '__cached__',
 '__doc__',
 '__file__',
 '__loader__',
 '__name__',
 '__package__',
 '__path__',
 '__spec__',
 'carshare',
 'election',
 'election_geojson',
 'experiment',
 'gapminder',
 'iris',
 'medals_long',
 'medals_wide',
 'stocks',
 'tips',
 'wind']

```

```

import seaborn as sns
dir(sns)

```

```

['FacetGrid',
 'JointGrid',
 'PairGrid',
 '__builtins__',
 '__cached__',
 '__doc__',
 '__file__',
 '__loader__',
 '__name__',
 '__package__',
 '__path__',

```

```
titanic.head()
```

```
penguins.head()
```

```
tips.head()
```

	total_bill	tip	sex	smoker	day	time	size
0	16.99	1.01	Female	No	Sun	Dinner	2
1	10.34	1.66	Male	No	Sun	Dinner	3
2	21.01	3.50	Male	No	Sun	Dinner	3
3	23.68	3.31	Male	No	Sun	Dinner	2
4	24.59	3.61	Female	No	Sun	Dinner	4

titanic.tail()

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class	who	adult_male	deck	embark_town	alive	alone
886	0	2	male	27.0	0	0	13.00	S	Second	man	True	NaN	Southampton	no	True
887	1	1	female	19.0	0	0	30.00	S	First	woman	False	B	Southampton	yes	True
888	0	3	female	NaN	1	2	23.45	S	Third	woman	False	NaN	Southampton	no	False
889	1	1	male	26.0	0	0	30.00	C	First	man	True	C	Cherbourg	yes	True
890	0	3	male	32.0	0	0	7.75	Q	Third	man	True	NaN	Queenstown	no	True

tips.tail()

	total_bill	tip	sex	smoker	day	time	size
239	29.03	5.92	Male	No	Sat	Dinner	3
240	27.18	2.00	Female	Yes	Sat	Dinner	2
241	22.67	2.00	Male	Yes	Sat	Dinner	2
242	17.82	1.75	Male	No	Sat	Dinner	2
243	18.78	3.00	Female	No	Thur	Dinner	2

titanic.sample(5)

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class	who	adult_male	deck	embark_town	alive	alone
199	0	2	female	24.0	0	0	13.0000	S	Second	woman	False	NaN	Southampton	no	T
713	0	3	male	29.0	0	0	9.4833	S	Third	man	True	NaN	Southampton	no	T
354	0	3	male	NaN	0	0	7.2250	C	Third	man	True	NaN	Cherbourg	no	T
231	0	3	male	29.0	0	0	7.7750	S	Third	man	True	NaN	Southampton	no	T
583	0	1	male	36.0	0	0	40.1250	C	First	man	True	A	Cherbourg	no	T

```
print("shape(rows,columns):",titanic.shape)
print("shape(rows,columns):",tips.shape)
print("Size(total Values):",titanic.size)
print("Dimensions (ndim):",titanic.ndim)
```

```
shape(rows,columns): (891, 15)
shape(rows,columns): (244, 7)
Size(total Values): 13365
Dimensions (ndim): 2
```

```
print("Titanic Columns:\n", list(titanic.columns))
print("Tips Columns:\n", list(tips.columns))
print("\n Index:",titanic.index)
```

```
Titanic Columns:
['survived', 'pclass', 'sex', 'age', 'sibsp', 'parch', 'fare', 'embarked', 'class', 'who', 'adult_male', 'deck', 'embark_town', 'alive', 'alone']
Tips Columns:
['total_bill', 'tip', 'sex', 'smoker', 'day', 'time', 'size']

Index: RangeIndex(start=0, stop=891, step=1)
```

titanic.dtypes

```

0
survived    int64
pclass      int64
sex         object
age         float64
sibsp       int64
parch       int64
fare        float64
embarked    object
class       category
who         object
adult_male   bool
deck        category
embark_town  object
alive       object
alone       bool

```

dtype: object

titanic.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 15 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   survived    891 non-null    int64
1   pclass      891 non-null    int64
2   sex         891 non-null    object
3   age         714 non-null    float64
4   sibsp       891 non-null    int64
5   parch       891 non-null    int64
6   fare        891 non-null    float64
7   embarked    889 non-null    object
8   class       891 non-null    category
9   who         891 non-null    object
10  adult_male   891 non-null    bool
11  deck        203 non-null    category
12  embark_town  889 non-null    object
13  alive       891 non-null    object
14  alone       891 non-null    bool
dtypes: bool(2), category(2), float64(2), int64(4), object(5)
memory usage: 80.7+ KB

```

tips.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 244 entries, 0 to 243
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   total_bill  244 non-null    float64
1   tip         244 non-null    float64
2   sex         244 non-null    category
3   smoker      244 non-null    category
4   day         244 non-null    category
5   time        244 non-null    category
6   size        244 non-null    int64
dtypes: category(4), float64(2), int64(1)
memory usage: 7.4 KB

```

titanic.describe()

	survived	pclass	age	sibsp	parch	fare
count	891.000000	891.000000	714.000000	891.000000	891.000000	891.000000
mean	0.383838	2.308642	29.699118	0.523008	0.381594	32.204208
std	0.486592	0.836071	14.526497	1.102743	0.806057	49.693429
min	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000
25%	0.000000	2.000000	20.125000	0.000000	0.000000	7.910400
50%	0.000000	3.000000	28.000000	0.000000	0.000000	14.454200
75%	1.000000	3.000000	38.000000	1.000000	0.000000	31.000000
max	1.000000	3.000000	80.000000	8.000000	6.000000	512.329200

```
tips.describe()
```

	total_bill	tip	size
count	244.000000	244.000000	244.000000
mean	19.785943	2.998279	2.569672
std	8.902412	1.383638	0.951100
min	3.070000	1.000000	1.000000
25%	13.347500	2.000000	2.000000
50%	17.795000	2.900000	2.000000
75%	24.127500	3.562500	3.000000
max	50.810000	10.000000	6.000000

```
titanic.describe(include="all")
```

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class	who	adult_male	deck	...
count	891.000000	891.000000	891	714.000000	891.000000	891.000000	891.000000	889	891	891	891	203	...
unique	NaN	NaN	2	NaN	NaN	NaN	NaN	3	3	3	2	7	...
top	NaN	NaN	male	NaN	NaN	NaN	NaN	S	Third	man	True	C	...
freq	NaN	NaN	577	NaN	NaN	NaN	NaN	644	491	537	537	59	...
mean	0.383838	2.308642	NaN	29.699118	0.523008	0.381594	32.204208	NaN	NaN	NaN	NaN	NaN	...
std	0.486592	0.836071	NaN	14.526497	1.102743	0.806057	49.693429	NaN	NaN	NaN	NaN	NaN	...
min	0.000000	1.000000	NaN	0.420000	0.000000	0.000000	0.000000	NaN	NaN	NaN	NaN	NaN	...
25%	0.000000	2.000000	NaN	20.125000	0.000000	0.000000	7.910400	NaN	NaN	NaN	NaN	NaN	...
50%	0.000000	3.000000	NaN	28.000000	0.000000	0.000000	14.454200	NaN	NaN	NaN	NaN	NaN	...
75%	1.000000	3.000000	NaN	38.000000	1.000000	0.000000	31.000000	NaN	NaN	NaN	NaN	NaN	...
max	1.000000	3.000000	NaN	80.000000	8.000000	6.000000	512.329200	NaN	NaN	NaN	NaN	NaN	...

```
age = titanic["age"].dropna() # drop missing values just for this calculation
```

```
print("Mean          :", round(age.mean(), 2))
print("Median        :", round(age.median(), 2))
print("Mode          :", age.mode()[0])
print("Variance       :", round(age.var(), 2))
print("Standard Deviation :", round(age.std(), 2))
print("Q1 (25th percentile):", age.quantile(0.25))
print("Q3 (75th percentile):", age.quantile(0.75))
print("90th Percentile   :", age.quantile(0.90))
print("Range            :", age.max() - age.min())
print("IQR (Q3 - Q1)     :", age.quantile(0.75) - age.quantile(0.25))
```

```
Mean          : 29.7
Median        : 28.0
Mode          : 24.0
Variance       : 211.02
Standard Deviation : 14.53
Q1 (25th percentile): 20.125
Q3 (75th percentile): 38.0
90th Percentile : 50.0
Range         : 79.58
IQR (Q3 - Q1) : 17.875
```


```
missing_count = titanic.isnull().sum()
missing_percent = (titanic.isnull().sum() / len(titanic) * 100).round(2)

missing_summary = pd.DataFrame({
    "Missing Count": missing_count,
    "Missing %": missing_percent
}).sort_values("Missing %", ascending=False)


missing_summary[missing_summary["Missing Count"] > 0]
```

	Missing Count	Missing %
deck	688	77.22
age	177	19.87
embarked	2	0.22
embark_town	2	0.22


```
# Visualization 1: Heatmap of missing values
plt.figure(figsize=(9, 5))
sns.heatmap(titanic.isnull(), cbar=False, cmap="viridis")
plt.title("Missing Value Heatmap - Titanic Dataset")
plt.show()
```



```
# Visualization 1: Heatmap of missing values
plt.figure(figsize=(9, 5))
sns.heatmap(tips.isnull(), cbar=False, cmap="viridis")
plt.title("Missing Value Heatmap - Tips Dataset")
plt.show()
```



```
# Visualization 2: Bar chart of missing value percentage
plt.figure(figsize=(8, 4))
missing_summary[missing_summary["Missing Count"] > 0]["Missing %"].plot(kind="bar", color="salmon")
plt.ylabel("Missing %")
plt.title("Percentage of Missing Values per Column")
plt.xticks(rotation=45)
plt.show()
```



```
df_demo = titanic.copy() # always work on a COPY, never the original!

# 1. Drop rows with any missing value (use sparingly – can lose a lot of data)
df_drop = df_demo.dropna()
print("After dropna():", df_drop.shape)

# 2. Fill numerical column with MEAN
df_demo["age_mean_filled"] = df_demo["age"].fillna(df_demo["age"].mean())

# 3. Fill numerical column with MEDIAN (better when data is skewed / has outliers)
df_demo["age_median_filled"] = df_demo["age"].fillna(df_demo["age"].median())

# 4. Fill categorical column with MODE (most frequent value)
df_demo["embarked_mode_filled"] = df_demo["embarked"].fillna(df_demo["embarked"].mode()[0])

# 5. Forward Fill – carry the previous valid value forward
df_demo["age_ffill"] = df_demo["age"].ffill()

# 6. Backward Fill – carry the next valid value backward
df_demo["age_bfill"] = df_demo["age"].bfill()

# 7. Interpolation – estimate missing values using surrounding data
df_demo["age_interpolated"] = df_demo["age"].interpolate()
```



```
# 8. Fill with a CONSTANT value
df_demo["deck_constant_filled"] = df_demo["deck"].astype("object").fillna("Unknown")

df_demo[["age", "age_mean_filled", "age_median_filled",
         "age_ffill", "age_bfill", "age_interpolated"]].head(8)
```

After dropna(): (182, 15)

	age	age_mean_filled	age_median_filled	age_ffill	age_bfill	age_interpolated
0	22.0	22.000000	22.0	22.0	22.0	22.0
1	38.0	38.000000	38.0	38.0	38.0	38.0
2	26.0	26.000000	26.0	26.0	26.0	26.0
3	35.0	35.000000	35.0	35.0	35.0	35.0
4	35.0	35.000000	35.0	35.0	35.0	35.0
5	NaN	29.699118	28.0	35.0	54.0	44.5
6	54.0	54.000000	54.0	54.0	54.0	54.0
7	2.0	2.000000	2.0	2.0	2.0	2.0

```
print("Number of duplicate rows in Titanic:", titanic.duplicated().sum())

# Let's create an artificial example to see duplicates in action
demo = titanic.head(5).copy()
demo = pd.concat([demo, demo.iloc[[0, 2]]], ignore_index=True) # manually duplicate 2 rows
print("\nDemo dataset shape (with duplicates):", demo.shape)
print("Duplicate rows found:", demo.duplicated().sum())

demo[demo.duplicated(keep=False)] # show all copies of duplicated rows
```

Number of duplicate rows in Titanic: 107

Demo dataset shape (with duplicates): (7, 15)
Duplicate rows found: 2

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class	who	adult_male	deck	embark_town	alive	alone
0	0	3	male	22.0	1	0	7.250	S	Third	man	True	NaN	Southampton	no	False
2	1	3	female	26.0	0	0	7.925	S	Third	woman	False	NaN	Southampton	yes	True
5	0	3	male	22.0	1	0	7.250	S	Third	man	True	NaN	Southampton	no	False
6	1	3	female	26.0	0	0	7.925	S	Third	woman	False	NaN	Southampton	yes	True

```
demo_cleaned = demo.drop_duplicates()
print("Shape before:", demo.shape, " | Shape after drop_duplicates():", demo_cleaned.shape)
```

Shape before: (7, 15) | Shape after drop_duplicates(): (5, 15)

```
print("Unique classes in 'pclass':", titanic["pclass"].unique())
print("Number of unique classes :", titanic["pclass"].nunique())

print("\nValue counts for 'embarked':")
print(titanic["embarked"].value_counts())

print("\nValue counts (as %) for 'sex':")
print(titanic["sex"].value_counts(normalize=True).round(3) * 100)
```

Unique classes in 'pclass': [3 1 2]
Number of unique classes : 3


Value counts for 'embarked':
embarked
S 644
C 168
Q 77
Name: count, dtype: int64

Value counts (as %) for 'sex':
sex
male 64.8
female 35.2
Name: proportion, dtype: float64


```
numeric_titanic = titanic.select_dtypes(include=np.number)
corr_matrix = numeric_titanic.corr()
corr_matrix
```

	survived	pclass	age	sibsp	parch	fare
survived	1.000000	-0.338481	-0.077221	-0.035322	0.081629	0.257307
pclass	-0.338481	1.000000	-0.369226	0.083081	0.018443	-0.549500
age	-0.077221	-0.369226	1.000000	-0.308247	-0.189119	0.096067
sibsp	-0.035322	0.083081	-0.308247	1.000000	0.414838	0.159651
parch	0.081629	0.018443	-0.189119	0.414838	1.000000	0.216225
fare	0.257307	-0.549500	0.096067	0.159651	0.216225	1.000000

```
plt.figure(figsize=(8, 6))
sns.heatmap(corr_matrix, annot=True, fmt=".2f", cmap="coolwarm", center=0)
plt.title("Correlation Heatmap – Titanic (numeric columns)")
plt.show()
```



```
plt.figure(figsize=(7, 4))
sns.boxplot(x=titanic["fare"], color="orange")
plt.title("Boxplot of Fare – Outliers appear as individual dots")
plt.show()
```



```
Q1 = titanic["fare"].quantile(0.25)
Q3 = titanic["fare"].quantile(0.75)
IQR = Q3 - Q1
```

```
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR
```

```

outliers = titanic[(titanic["fare"] < lower_bound) | (titanic["fare"] > upper_bound)]
print(f"Lower bound: {lower_bound:.2f} Upper bound: {upper_bound:.2f}")
print("Number of outliers in 'fare':", outliers.shape[0])

```

```

Lower bound: -26.72 Upper bound: 65.63
Number of outliers in 'fare': 116

```

```

from scipy import stats

```

```

fare_no_na = titanic["fare"].dropna()
z_scores = np.abs(stats.zscore(fare_no_na))

```

```

z_outliers = fare_no_na[z_scores > 3] # points more than 3 std-devs from the mean
print("Number of outliers using Z-score (|z| > 3):", z_outliers.shape[0])

```

```

Number of outliers using Z-score (|z| > 3): 20

```

```

titanic_no_outliers = titanic[(titanic["fare"] >= lower_bound) & (titanic["fare"] <= upper_bound)]
print("Original shape :", titanic.shape)
print("After removing fare outliers:", titanic_no_outliers.shape)

```

```

Original shape : (891, 15)
After removing fare outliers: (775, 15)

```

```

housing_num = housing.select_dtypes(include=np.number).dropna()
sample_cols = housing_num.columns[:2] # take 2 numeric columns to compare visually

```

```


scalers = {
    "Original": housing_num[sample_cols],
    "StandardScaler": pd.DataFrame(StandardScaler().fit_transform(housing_num[sample_cols]), columns=sample_cols),
    "MinMaxScaler": pd.DataFrame(MinMaxScaler().fit_transform(housing_num[sample_cols]), columns=sample_cols),
    "RobustScaler": pd.DataFrame(RobustScaler().fit_transform(housing_num[sample_cols]), columns=sample_cols),
    "Normalizer": pd.DataFrame(Normalizer().fit_transform(housing_num[sample_cols]), columns=sample_cols),
}

```

```

fig, axes = plt.subplots(1, 5, figsize=(22, 4))
for ax, (name, data) in zip(axes, scalers.items()):
    sns.histplot(data[sample_cols[0]], ax=ax, kde=True, color="teal")
    ax.set_title(name)
plt.suptitle(f"Effect of Different Scalers on '{sample_cols[0]}'")
plt.tight_layout()
plt.show()

```



```

demo_enc = titanic[["sex", "embarked", "pclass"]].dropna().copy()

```

```

# 1. Label Encoding - assigns each category an integer (0, 1, 2, ...)
le = LabelEncoder()
demo_enc["sex_label"] = le.fit_transform(demo_enc["sex"])

```

```

# 2. One-Hot Encoding - creates a separate 0/1 column for every category
one_hot = pd.get_dummies(demo_enc["embarked"], prefix="embarked")

```

```

# 3. Ordinal Encoding - like Label Encoding, but you CONTROL the order
# (pclass 1 = best, 3 = worst, so we keep that natural order)
ordinal = OrdinalEncoder(categories=[[1, 2, 3]])
demo_enc["pclass_ordinal"] = ordinal.fit_transform(demo_enc[["pclass"]])

```


```

result = pd.concat([demo_enc[["sex", "sex_label"]], one_hot,
                    demo_enc[["pclass", "pclass_ordinal"]]], axis=1)
result.head()


```

	sex	sex_label	embarked_C	embarked_Q	embarked_S	pclass	pclass_ordinal
0	male	1	False	False	True	3	2.0
1	female	0	True	False	False	1	0.0
2	female	0	False	False	True	3	2.0
3	female	0	False	False	True	1	0.0
4	male	1	False	False	True	3	2.0


```
plt.figure(figsize=(7, 4))
sns.histplot(titanic["age"], bins=30, kde=True, color="steelblue")
plt.title("Histogram of Passenger Age (Titanic)")
plt.xlabel("Age")
plt.show()
```




```
plt.figure(figsize=(6, 4))
sns.countplot(x="pclass", data=titanic, hue="pclass", palette="Set2", legend=False)
plt.title("Passenger Count by Class")
plt.show()
```




```
plt.figure(figsize=(6, 4))
sns.barplot(x="pclass", y="survived", data=titanic, hue="pclass", palette="Set1", legend=False)
plt.title("Average Survival Rate by Class")
plt.ylabel("Survival Rate")
plt.show()
```




```
plt.figure(figsize=(7, 4))
sns.violinplot(x="sex", y="age", data=titanic, hue="sex", palette="muted", legend=False)
plt.title("Age Distribution by Gender (Violin Plot)")
plt.show()
```




```
plt.figure(figsize=(7, 4))
sns.scatterplot(x="sepal_length", y="petal_length", hue="species", data=iris, palette="deep")
plt.title("Sepal Length vs Petal Length (Iris)")
plt.show()
```



```
sns.pairplot(iris, hue="species", palette="husl", diag_kind="kde")
plt.suptitle("Pairplot of Iris Dataset", y=1.02)
plt.show()
```




```
plt.figure(figsize=(6, 5))
sns.heatmap(iris.drop(columns="species").corr(), annot=True, cmap="YlGnBu")
plt.title("Correlation Heatmap - Iris Dataset")
plt.show()
```




```
bill_by_size = tips.groupby("size")["total_bill"].mean().reset_index()

plt.figure(figsize=(7, 4))
sns.lineplot(x="size", y="total_bill", data=bill_by_size, marker="o", color="darkgreen")
plt.title("Average Total Bill by Party Size (Tips Dataset)")
plt.xlabel("Party Size")
plt.ylabel("Average Total Bill ($)")
plt.show()
```




```
plt.figure(figsize=(5, 5))
titanic["embarked"].value_counts().plot(
    kind="pie", autopct="%1.1f%%", colors=sns.color_palette("Set3"), startangle=90)
plt.title("Passengers by Embarkation Port")
plt.ylabel("")
plt.show()
```


Passengers by Embarkation Port




```
plt.figure(figsize=(7, 4))
sns.kdeplot(data=titanic, x="fare", hue="survived", fill=True, common_norm=False, palette="Set1")
plt.title("Fare Distribution by Survival (KDE Plot)")
plt.xlim(0, 300)
plt.show()
```




```
plt.figure(figsize=(7, 4))
sns.stripplot(x="day", y="total_bill", data=tips, hue="day", jitter=True, palette="cool", legend=False)
plt.title("Total Bill by Day (Strip Plot)")
plt.show()
```





```
plt.figure(figsize=(7, 4))
sns.swarmplot(x="day", y="total_bill", data=tips, hue="day", palette="autumn", size=4, legend=False)
plt.title("Total Bill by Day (Swarm Plot)")
plt.show()
```




```
sns.jointplot(x="total_bill", y="tip", data=tips, kind="reg", color="purple", height=6)
plt.suptitle("Total Bill vs Tip (Jointplot with Regression Line)", y=1.02)
plt.show()
```



```
plt.figure(figsize=(7, 3))
sns.kdeplot(tips["tip"], color="teal")
sns.rugplot(tips["tip"], color="black")
plt.title("Tip Amount - KDE with Rug Plot")
plt.show()
```



```
plt.figure(figsize=(7, 4))
sns.ecdfplot(data=titanic, x="fare", hue="survived", palette="Dark2")
plt.title("ECDF of Fare by Survival")
plt.show()
```



```
# groupby() + single aggregation
tips.groupby("day", observed=True)["total_bill"].mean()
```

```
total_bill
day
Thur    17.682742
Fri     17.151579
Sat     20.441379
Sun     21.410000
```

```
dtype: float64
```

```
# groupby() + agg() with multiple statistics at once
tips.groupby("day", observed=True)["total_bill"].agg(["mean", "median", "std", "count"])
```

```
mean median      std count
day
Thur    17.682742   16.20  7.886170    62
Fri     17.151579   15.38  8.302660    19
Sat     20.441379   18.24  9.480419    87
Sun     21.410000   19.63  8.832122    76
```

```
# groupby() with multiple columns
tips.groupby(["day", "time", observed=True]["tip"].mean().unstack()
```

time	Lunch	Dinner
day		
Thur	2.767705	3.000000
Fri	2.382857	2.940000
Sat	NaN	2.993103
Sun	NaN	3.255132

```
# pivot_table() - similar to groupby, but in a spreadsheet-like format
pd.pivot_table(tips, values="total_bill", index="day", columns="time",
               aggfunc="mean", observed=True)
```

time	Lunch	Dinner
day		
Thur	17.664754	18.780000
Fri	12.845714	19.663333
Sat	NaN	20.441379
Sun	NaN	21.410000

```
# crosstab() - frequency table between two categorical columns
pd.crosstab(titanic["pclass"], titanic["survived"], margins=True)
```

survived	0	1	All
pclass			
1	80	136	216
2	97	87	184
3	372	119	491
All	549	342	891

```
# Start fresh with a clean copy
clean_df = titanic.copy()
print("Starting shape:", clean_df.shape)
clean_df.head()
```

Starting shape: (891, 15)

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class	who	adult_male	deck	embark_town	alive	alone
0	0	3	male	22.0	1	0	7.2500	S	Third	man	True	NaN	Southampton	no	False
1	1	1	female	38.0	1	0	71.2833	C	First	woman	False	C	Cherbourg	yes	False
2	1	3	female	26.0	0	0	7.9250	S	Third	woman	False	NaN	Southampton	yes	True
3	1	1	female	35.0	1	0	53.1000	S	First	woman	False	C	Southampton	yes	False
4	0	3	male	35.0	0	0	8.0500	S	Third	man	True	NaN	Southampton	no	True

```
# 'age' -> fill with median (skewed, has outliers)
clean_df["age"] = clean_df["age"].fillna(clean_df["age"].median())

# 'embarked' -> fill with mode (categorical, very few missing)
clean_df["embarked"] = clean_df["embarked"].fillna(clean_df["embarked"].mode()[0])

# 'deck' -> too many missing (~77%) -> drop the column entirely
clean_df = clean_df.drop(columns=["deck"])

# drop rows still missing anything critical (safety net)
clean_df = clean_df.dropna(subset=["embark_town"])

print("Missing values remaining:\n", clean_df.isnull().sum()[clean_df.isnull().sum() > 0])
```

Missing values remaining:
Series([], dtype: int64)

```
before = clean_df.shape[0]
clean_df = clean_df.drop_duplicates()
after = clean_df.shape[0]
print(f"Removed {before - after} duplicate rows. New shape: {clean_df.shape}")
```

Removed 116 duplicate rows. New shape: (773, 14)

```
Q1, Q3 = clean_df["fare"].quantile([0.25, 0.75])
IQR = Q3 - Q1
lower, upper = Q1 - 1.5 * IQR, Q3 + 1.5 * IQR

n_outliers = clean_df[(clean_df["fare"] < lower) | (clean_df["fare"] > upper)].shape[0]
print(f"Fare outliers detected: {n_outliers}")

# Cap (winsorize) rather than delete, to avoid losing too many rows
clean_df["fare"] = clean_df["fare"].clip(lower=lower, upper=upper)
print("Fare range after capping:", clean_df["fare"].min(), "-", clean_df["fare"].max())
```

Fare outliers detected: 100
Fare range after capping: 0.0 - 71.675

```
# Drop redundant / high-cardinality / identifier-like columns first
clean_df = clean_df.drop(columns=["who", "adult_male", "alive", "class", "embark_town"])

# One-Hot Encode nominal columns
clean_df = pd.get_dummies(clean_df, columns=["sex", "embarked"], drop_first=True)

clean_df.head()
```

	survived	pclass	age	sibsp	parch	fare	alone	sex_male	embarked_Q	embarked_S
0	0	3	22.0	1	0	7.2500	False	True	False	True
1	1	1	38.0	1	0	71.2833	False	False	False	False
2	1	3	26.0	0	0	7.9250	True	False	False	True
3	1	1	35.0	1	0	53.1000	False	False	False	True
4	0	3	35.0	0	0	8.0500	True	True	False	True

```
num_cols = ["age", "fare"]
scaler = StandardScaler()
clean_df[num_cols] = scaler.fit_transform(clean_df[num_cols])

clean_df[num_cols].describe().round(2)
```

	age	fare
count	773.00	773.00
mean	0.00	-0.00
std	1.00	1.00
min	-2.12	-1.18
25%	-0.62	-0.82
50%	-0.11	-0.46
75%	0.47	0.33
max	3.68	2.04

```
print("Final cleaned shape:", clean_df.shape)
print("\nRemaining missing values:", clean_df.isnull().sum().sum())
clean_df.head()
```

Final cleaned shape: (773, 10)

Remaining missing values: 0

	survived	pclass	age	sibsp	parch	fare	alone	sex_male	embarked_Q	embarked_S
0	0	3	-0.548619	1	0	-0.851827	False	True	False	True
1	1	1	0.617360	1	0	2.019543	False	False	False	False
2	1	3	-0.257124	0	0	-0.821559	True	False	False	True
3	1	1	0.398739	1	0	1.204170	False	False	False	True
4	0	3	0.398739	0	0	-0.815954	True	True	False	True

```
# Save the final cleaned dataset to CSV
clean_df.to_csv("titanic_cleaned.csv", index=False)
print(" Saved as 'titanic_cleaned.csv'")
```

Saved as 'titanic_cleaned.csv'

```
# STEP 1 (example, already solved): First look at the data
print(tips.shape)
tips.head()
```

(244, 7)

	total_bill	tip	sex	smoker	day	time	size
0	16.99	1.01	Female	No	Sun	Dinner	2
1	10.34	1.66	Male	No	Sun	Dinner	3
2	21.01	3.50	Male	No	Sun	Dinner	3
3	23.68	3.31	Male	No	Sun	Dinner	2
4	24.59	3.61	Female	No	Sun	Dinner	4

```
# STEP 2: Check for missing values
# TODO: use isnull().sum() on the tips dataset
tips.isnull().sum()
```

	0
total_bill	0
tip	0
sex	0
smoker	0
day	0
time	0
size	0

dtype: int64


```
# STEP 3: Check for duplicate rows
# TODO: use duplicated().sum()
tips.duplicated().sum()
```

np.int64(1)

```
# STEP 4: Get summary statistics
# TODO: use describe()
tips.describe()
```

	total_bill	tip	size
count	244.000000	244.000000	244.000000
mean	19.785943	2.998279	2.569672
std	8.902412	1.383638	0.951100
min	3.070000	1.000000	1.000000
25%	13.347500	2.000000	2.000000
50%	17.795000	2.900000	2.000000
75%	24.127500	3.562500	3.000000
max	50.810000	10.000000	6.000000

```
# STEP 5: Visualize the relationship between total_bill and tip
# TODO: create a scatter plot, colored by 'time'
plt.figure(figsize=(7, 4))
sns.scatterplot(x="total_bill", y="tip", hue="time", data=tips, palette="Set1")
plt.title("Total Bill vs Tip, colored by Time of Day")
plt.show()
```



```
pd.read_csv("titanic_cleaned.csv")
```

	survived	pclass	age	sibsp	parch	fare	alone	sex_male	embarked_Q	embarked_S
0	0	3	-0.548619	1	0	-0.851827	False	True	False	True
1	1	1	0.617360	1	0	2.019543	False	False	False	False
2	1	3	-0.257124	0	0	-0.821559	True	False	False	True
3	1	1	0.398739	1	0	1.204170	False	False	False	True
4	0	3	0.398739	0	0	-0.815954	True	True	False	True
...	...	...	...	...	...	...	...	...	...	...
768	0	3	0.690233	0	5	0.129088	False	False	True	False
769	1	1	-0.767240	0	0	0.168324	True	False	False	True
770	0	3	-0.111377	1	2	-0.125390	False	False	False	True
771	1	1	-0.257124	0	0	0.168324	True	True	False	False
772	0	3	0.180118	0	0	-0.829407	True	True	True	False

773 rows × 10 columns

```
# Exercise 1: Load Iris and print the first five rows
# TODO: complete this
iris_ex = sns.load_dataset("iris")
print(iris_ex.head())
```

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	setosa
1	4.9	3.0	1.4	0.2	setosa
2	4.7	3.2	1.3	0.2	setosa
3	4.6	3.1	1.5	0.2	setosa
4	5.0	3.6	1.4	0.2	setosa

```
# Exercise 2: Find missing values in the Titanic dataset
# TODO: complete this
print(titanic.isnull().sum())
```


```
survived      0
pclass        0
sex            0
age           177
sibsp          0
parch          0
fare           0
embarked       2
class          0
who            0
adult_male     0
deck          688
embark_town    2
alive          0
alone          0
dtype: int64
```

```
# Exercise 3: Remove duplicate rows from the Tips dataset (if any)
# TODO: complete this
```


```
tips_no_dup = tips.drop_duplicates()
print("Before:", tips.shape, "After:", tips_no_dup.shape)
```

Before: (244, 7) After: (243, 7)

```
# Exercise 4: Plot a histogram of the 'total_bill' column from Tips
# TODO: complete this
plt.figure(figsize=(6, 4))
sns.histplot(tips["total_bill"], bins=20, kde=True, color="coral")
plt.title("Histogram of Total Bill")
plt.show()
```



```
# Exercise 5: Create a correlation heatmap for the Penguins dataset
# TODO: complete this
plt.figure(figsize=(6, 5))
sns.heatmap(penguins.select_dtypes(include=np.number).corr(), annot=True, cmap="coolwarm")
plt.title("Penguins Correlation Heatmap")
plt.show()
```



Mini Assignment on Penguins Dataset

```
#Number of rows and columns
print("shape(rows,columns):", penguins.shape)
```

shape(rows,columns): (344, 7)

```
# Number of missing values (per column)
missing_count = penguins.isnull().sum()
missing_percent = (penguins.isnull().sum() / len([penguins])) * 100).round(2)

missing_summary = pd.DataFrame({
    "Missing Count": missing_count,
    "Missing %": missing_percent
}).sort_values("Missing %", ascending=False)

missing_summary[missing_summary["Missing Count"] > 0]
```

	Missing Count	Missing %
sex	11	1100.0
bill_depth_mm	2	200.0
bill_length_mm	2	200.0
flipper_length_mm	2	200.0
body_mass_g	2	200.0

```
#Number of duplicate rows
print("Number of duplicate rows in Penguins:", penguins.duplicated().sum())
```

Number of duplicate rows in Penguins: 0

```
#A correlation matrix / heatmap of the numeric columns
numeric_penguins = penguins.select_dtypes(include=np.number)
corr_matrix = numeric_penguins.corr()
corr_matrix
```


	bill_length_mm	bill_depth_mm	flipper_length_mm	body_mass_g
bill_length_mm	1.000000	-0.235053	0.656181	0.595110
bill_depth_mm	-0.235053	1.000000	-0.583851	-0.471916
flipper_length_mm	0.656181	-0.583851	1.000000	0.871202
body_mass_g	0.595110	-0.471916	0.871202	1.000000

```
#5. A boxplot showing `bill_length_mm` grouped by `species`
plt.figure(figsize=(8, 6))

sns.boxplot(
    x="species",
    y="bill_length_mm",
    data=penguins
)

plt.title("Bill Length by Penguin Species")
plt.xlabel("Species")
plt.ylabel("Bill Length (mm)")

plt.show()
```

```
# 6. A pairplot colored by `species`  
sns.pairplot(penguins, hue="species", palette="Set2", diag_kind="kde")  
plt.suptitle("Pairplot of Penguins Dataset", y=1.02)  
plt.show()
```

Pairplot of Penguins Dataset



```
sns.pairplot(  
    data=penguins,  
    hue="species",  
    diag_kind="hist"  
)
```

```
plt.show()
```

